

# Extracted the style from the image with Deep Neural Network

Bingzhe Wu

School of Mathematical Science

Peking University

Email:wubingzhe@pku.edu.cn

**Abstract**—In the field of visual art, especially painting, humans have mastered the skill to create unique visual experience through composing a complex interplay between the content and style of an image. In other area of computer vision such as object detection and recognition, Convolutional neural networks have recently enjoyed a great success in large-scale image recognition[1] which has become possible due to large public image repositories, such as ImageNet(Deng et al., 2009), and high-performance computing systems, such as GPUs or large-scale distributed clusters(Dean et al. 2012). In this project, we use a neural representations to separate and recombine content and style of arbitrary images. This work also offers a algorithmic understanding of how humans create and perceive artistic imagery.

## I. INTRODUCTION

In this project, we use a most powerful Deep Neural Networks in image processing called Convolutional Neural Networks. When Convolutional Neural Networks are trained on object recognition, they develop a representation of the image that makes object information increasingly explicit along the processing hierarchy. Therefore, along the processing hierarchy of the network, the input image is transformed into representations that care about the actual content of the image compared to its detailed pixel. In terms of this, we can directly visualise the information each layer contains about the input image by reconstructing the image only from the feature maps in that layer. High layers in the network capture the high-level content in terms of objects and their arrangement in input image but do not constrain the exact pixel values of reconstruction. In contrast, reconstruction from the lower layers simply reproduce the exact pixel values of the original image. In this project, we refer to the feature responses in high layers of the networks as the content representation.

To obtain a representation of the style of an input image, we use a feature space originally designed to capture texture information. This feature space we used is built on top of the filter responses in each layer of the network. It consists of the correlations between the different filter responses over the spatial extent of the feature maps. (See Implementation for details). By including the feature correlations of multiple layers, we obtain a multi-scale representation of the input image, which captures its style information but not the details of content and its arrangement.

Again, we can visualise the information captured by these style feature spaces built on different layers of the network by

constructing an image that matches the style representation of a given input image. Indeed reconstruction from the style features from the style features produce texturised versions of the input image that capture its general appearance in terms of color and localised structures. Moreover, the size and complexity of local image structures from the input image increases along the hierarchy, a result that can be explained by the increasing receptive field sizes and feature complexity. We refer to this multi-scale representation as style representation.

In this project, we can manipulate both representation independently to produce new meaningful image. To demonstrate this work, we generate image mix the content and style representation from two different source images. In particular,

The images are synthesised by finding an image that simultaneously matches the content representation of the photograph and the style representation of the respective piece of art. While the global arrangement of the original photograph is preserved, the colors and local structures that compose the global scenery are provided by the artwork. Effectively, this renders the photograph in the style of the artwork, such that the appearance of the synthesised image resembles the work of art, even though it shows the same content as photograph.

Of course, image content and style cannot be completely disentangled. When synthesising an image that combine the content of one image with the style of another image, there usually does not exist an image that perfectly matches both constraints at the same time. However, the loss function we minimise during image synthesis contains two terms for content and style respectively, that are well separated. We can therefore smoothly regulate the emphasis on either reconstructing the content or the style.

In my final project, we set up an artificial system that achieves a separation of image content from style, thus allowing to recast the content of one image in the style of any other image. We demonstrate this by creating new, artistic images that combine the style of several well-known paintings with the content of an arbitrarily chosen photograph that I take. In particular, we derive the neural representations for the content and style of an image from the feature responses of high-performing Deep Neural Networks trained on object recognition.

As outlined above,

The rest of the report is organized as follows: Section 2 provides a background for CNN and VGG model, and also

described the datasets for training model we used. Section 3 presents our described the method we used. Section 4 describes implementation details. Section 5 shows our experiment result.

## II. BACKGROUND

### A. DATASETS

ImageNet is a dataset of over 15 million labeled high-resolution images belong to roughly 22,000 categories. The images were collected from the web and labeled by human labelers using Amazon's Mechanical Turk crowd-sourcing tool. Starting in 2010, as part of the Pascal Visual Object Challenge, an annual competition called the ImageNet Large-Scale Visual Recognition Challenge (ILSVRC) has been held. ILSVRC uses a subset of ImageNet with roughly 1000 images in each of 1000 categories. In all, there are roughly 1.2 million training images, 50,000 validation images, and 150,000 testing images.

ImageNet consists of variable-resolution images, while our system requires a constant input dimensionality. In our training process, we down-sampled the images to a fixed resolution of  $256 \times 256$ . Given an image, we first rescaled the image such that the shorter side was of length 256, and then cropped out the central  $256 \times 256$  patch from the resulting image. We did not pre-process the images in any other way, except for subtracting the mean activity over the training set from each pixel. Then we use these pictures for training a VGG19 (see it in next subsection) model by Caffe (A deep learning framework).

### B. CNN Basics

Convolutional neural network (CNN) is first inspired by research in neuroscience. After over twenty years of convolution, CNN has been gaining more and more distinction in research fields, such as computer vision, pattern recognition, NLP. As a classical supervised learning algorithm, CNN employs a feedforward process for recognition and backward path for training.

A typical CNN is composed of two components: a feature extractor and a classifier. The feature extractor is used to filter input images into feature maps that represent various features of the image. These features may include corners, lines, circular arch, etc., which are relatively invariant to position shifting or distortions. The output of the feature extractor is a low-dimensional vector containing these features. This vector is then fed into the classifier, which is usually based on traditional artificial neural networks. The purpose of this classifier is to decide the likelihood of categories that the input image might belong to.

A typical CNN is composed of multiple computation layers. For example, the feature extractor may consist of several convolutional layers and optional sub-sampling layers. Figure 1 illustrates the computation of a convolutional layer. The convolutional layer receives  $N$  feature maps as input. Each input feature map is convolved by a shifting window of size  $S$ , which is normally smaller than  $K$ . A total of  $M$  output feature

maps will form the set of input feature maps for the next convolutional layer.




Fig. 1. Graph of a convolutional layer

### C. VGG Model

1) *ARCHITECTURE*: VGG model is a convolutional neural network that rivals human performance on a common visual object recognition benchmark task and was introduced and extensively described in [10].

In this project, we used VGG models for training. During training, the input to this ConvNets is a fixed-size  $224 \times 224$  RGB image. The only pre-processing we do is subtracting the mean RGB value, computed on the training set, from each pixel. The image is passed through a stack of convolution layers, where we use filters with a very small receptive field:  $3 \times 3$  (which is the smallest size to capture the notion of left/right, up/down, center). In one of the configurations we also utilize  $1 \times 1$  convolution filters, which can be seen as a linear transformation of the input channels (followed by non-linearity). The convolutional stride is fixed to 1 pixel; the spatial padding of conv. layer input is such that the spatial resolution is preserved after convolution, i.e. the padding is 1 pixel for  $3 \times 3$  conv. layers. Spatial pooling is carried out by five max-pooling layers, which follow some of the conv. layers (not all the conv. layers are followed by max-pooling). Max-pooling is performed over a  $2 \times 2$  pixel window, with stride 2.

A stack of convolutional layers (which has a different depth in different architectures) is followed by three Fully-Connected (FC) layers. However in our project, the fully-connected layers were removed.

All hidden layers are equipped with the rectification (ReLU) non-linearity. We note that none of our networks (except for one) contain Local Response Normalisation (LRN) normalisation.

The Convnet we used in the project are shown in Figure 2

## III. METHODS

The results presented in our project were generated on the basis of the VGG-Model (see details in section 2). We used the feature space provided by the 16 convolutional and 5 pooling layers of the 19 layer VGG-Network (see the architecture in Figure 2). We do not use any of the fully connected layers. The model is publicly available and can be explored in the Caffe framework. For image synthesis we found that replacing the max-pooling operation by average pooling improves the gradient flow and one obtains slightly more appealing results, which is why the images shown were generated with average pooling.



Fig. 2. Graph of a convolutional layer

Generally each layer in the network defines a non-linear filter bank whose complexity increases with the position of the layer in the network .Hence a given input image  $\vec{x}$  is encoded in each layer of the CNN by the filter responses to that image.A layer with  $N_l$  distinct filters has  $N_l$  feature maps each of size  $M_l$ ,where  $M_l$  is the height times the width of the feature maps. So the responses in a layer  $l$  can be stored in a matrix  $F^l \in R^{N_l \times M_l}$  where  $F_{ij}^l$  is the activation of the  $i^{th}$  filter at position  $j$  in layer  $l$